



tare

ZERO THE CONFIDENCE. WEIGH THE RECORD.

Brand kit for an agentic paper trader with a non-LLM Inspector.

Flight Recorder · Guarded vs Shadow · Attack Lab

METAPHOR

Tare is the empty weight you subtract so the scale reads what is actually there.

Stated confidence is gross. Calibrated probability is net.

The Inspector zeros the model before size is allowed to move.

Contents

01	Name & metaphor	Why Tare, not another LLM judge
02	Positioning	Brand spine, one-liner, what it is not
03	Audience	Judges, builders, operators
04	Logo system	Mark, wordmark, lockups, favicon
05	Clear space & misuse	Construction and prohibitions
06	Color	Semantic tokens only
07	Typography	Plex Sans / Plex Mono / scale
08	Voice	Vocabulary, taglines, microcopy
09	UI system	Radius, density, components
10	Product surfaces	Flight Recorder, Attack Lab, landing
11	Evidence rules	Gates, attacks, honesty aesthetic
12	Do / Don't	Implementation notes

Tare the model.

On a scale, tare is the mechanical act of zeroing out the empty container so the instrument reports true net weight. Tare the product does the same thing to an LLM trader: it subtracts inflated confidence and reads the calibrated probability underneath.

The name is short, unfashionable, and exact. It is a verb and a measurement. It does not celebrate intelligence. It describes a correction.

GROSS

Stated confidence

The model's P(TP before SL).
A claim. Often overweight.

TARE

Inspector correction

Historical hit rate,
independent sensors, hard
limits. The empty weight.

NET

Calibrated p / size

What is allowed to pass the
gate. The number that may
trade.

PRONUNCIATION & WRITING

SAY tair — same vowel as “care.” Not “tar-ee.”

WRITE tare wordmark · Tare in sentences · TARE only as a terminal stamp.

NEVER TareAI Tare Bot The Tare tare® ProsecutorAI

KEEP Inspector, Flight Recorder, Shadow, Gate G0–G2, Attack A1–A5.

ONE-LINER

An AI trader with brakes that tighten where the model can't be trusted.

Not a smarter judge. A scale.

A second AI reviewing the first shares the same blind spots. When the trader is fooled, the peer reviewer fails at the same moment. Tare never asks another model to grade a model. It asks the record, and independent evidence the trader cannot fake.

IT IS

- Rules-based market-structure setups
- LLM proposes take/skip + confidence
- Non-LLM Inspector sizes or vetoes
- Calibration against stated confidence
- Independent anomaly sensors
- Hard risk limits the model cannot edit
- Shadow book (unguarded counterfactual)
- Attack Lab that poisons trader inputs
- Flight Recorder of every decision

IT IS NOT

- A smarter LLM that reviews another LLM
- A signal group or call channel
- A meme-coin launcher
- A DeFi yield product
- A black-box “AI alpha” flex
- Web3 carnival branding
- Friendly toast notifications for harm
- Social trading, NFTs, or a mobile app
- Anything that hides a bad metric

Prove it in seconds.

Hackathon judges

Novelty, rigor, demo clarity

This is not another LLM trader. It has brakes. We attack them.

Agentic-trading builders

Trust, risk, adversarial robustness

Calibration vs stated confidence. Veto with numbers. Green vs red equity.

Technical reviewers

Honesty under attack

A4 slips past calibration-only. Full gate G2 catches it.

Live operator

Job-to-be-done


Show me whether the AI is overconfident — and whether the gate stopped it.

SUCCESS TEST


In 5 seconds: this product distrusts AI confidence. It does not celebrate it.

Mark, word, lockup.

Two parts, never mixed in tight chrome. The icon is the scale: white gross mass, red net mass, capsule beam, downward fulcrum. The wordmark is lowercase tare in Poppins Bold. No registered-mark glyph. No gavel, no brain, no chain link.



MARK · DARK



MARK · LIGHT

tare

WORDMARK · POPPINS BOLD · NO REGISTERED MARK



tare

PRIMARY LOCKUP

Clear space is the beam.

Clear space on all sides equals the height of the icon's capsule beam. Do not crowd the lockup. Minimum digital size: icon 28 px, wordmark 72 px wide. Favicon and app icon use the scale only.



CLEAR SPACE = 1× BEAM HEIGHT

- 01 **Left mass**
Stated confidence — gross, overweight. Always white/ink.
- 02 **Red mass**
The inflation we subtract. Only red allowed in the mark.
- 03 **Beam**
The comparison. Capsule, always level.
- 04 **Fulcrum**
The Inspector. Points down. Does not move with the model.
- 05 **Wordmark**
tare — Poppins Bold, lowercase. Never tare®.
- 06 **Icon tile**
Squirrelle app icon. Scale only — no word inside.

MISUSE

- Do not add “AI”, lightning, coins, gavels, chain links, or a ®.
- Do not recolor the red square. Do not paint the left mass green.
- Do not apply gradients, glows, drop shadows, or 3D bevels.
- Do not stretch, rotate past 0°, or place on busy photography.
- Do not set the wordmark in a serif, script, or condensed novelty face.
- Do not lock the mark to a cartoon mascot or scales-of-justice clipart.

Color is a verdict.

Near-monochrome canvas. Color is scarce and semantic. Green means guarded — the trade survived the gate. Red means shadow, veto, or harm. White is structure. There is no third accent family.

<code>bg.canvas</code> App / page background	<code>#0A0A0A</code>	<code>bg.panel</code> Cards, chart plot area	<code>#111111</code>
<code>bg.panel-elevated</code> Hover / selected rows	<code>#161616</code>	<code>border.subtle</code> Dividers, chart grid	<code>#222222</code>
<code>border.strong</code> Focus rings (non-danger)	<code>#333333</code>	<code>text.primary</code> Body	<code>#F0F0F0</code>
<code>text.muted</code> Labels, captions	<code>#A3A3A3</code>	<code>text.inverse</code> Text on light buttons only	<code>#0A0A0A</code>
<code>accent.guarded</code> Guarded equity, approve, held	<code>#22C55E</code>	<code>accent.danger</code> Shadow equity, veto, harm	<code>#EF4444</code>
<code>accent.danger-deep</code> Veto flash background	<code>#7F1D1D</code>	<code>accent.neutral</code> Chrome, gauges, bars	<code>#FFFFFF</code>

RULE · UI COLOR IS A VERDICT. EXCEPTION: THE ICON'S RED SQUARE IS PART OF THE MARK.

Grotesk authority.

Poppins Bold is the wordmark only. IBM Plex Sans carries UI and body. IBM Plex Mono is reserved for hashes, order IDs, veto numbers, and raw JSON. Tabular figures are mandatory for prices, percents, R-multiples, and confidence.

tare

Wordmark · Poppins Bold

Flight Recorder

Section 20–24

The Inspector sizes or vetoes using calibration.

Body 14–16

VETO · p_adj 0.29 < p_be 0.33 · anomaly 0.12

Meta / data 12–13

SCALE & RULES

- Headings: slight negative letter-spacing. No all-caps walls except stamps (VETO, G2, A4).
- Never mix a display serif into the product UI. Marketing may use Plex only.
- Confidence is always written as a probability that TP hits before SL.
- limits_hash is mono, truncated, treated as sealed — not decorative.

Expert witness.

Precise, calm, slightly severe. Short sentences. Numbers over adjectives. Willing to show bad results — honesty is a brand feature. Dry wit is allowed once. Hype is not.

TAGLINES

PRIMARY	Zero the confidence. Weigh the record.
ALTERNATE	Gross stated. Net calibrated.
PRODUCT	Brakes where the model can't be trusted.
ARGUMENT	Not another LLM judge. Evidence.
CHART	Guarded green. Unguarded red.

HOOKS THAT USE REAL NUMBERS

“This AI trader is 90% sure. Historically, when it says 90%, it’s right 41% of the time.”

“Another LLM won’t save you. Shared blind spots fail at the same moment.”

“We don’t ask a model to grade a model. We ask the record.”

PRODUCT VOCABULARY

Tare	Product name	Trader	LLM + setup layer
Inspector	Non-LLM gate — never “AI reviewer”	Flight Recorder	Dashboard / event log
Confidence	P(TP before SL), 0–100	Calibrated p	Empirical / Wilson-adjusted hit rate
Veto / Shrink / Approve	Hard verbs. No euphemism.	Shadow (unguarded)	Counterfactual book
G0 / G1 / G2	Unguarded / calibration / full	A1–A5	Named attacks

Stamped metal, not soft SaaS.

Radius	4 px default. 2–6 px maximum. Never pills on primary chrome.
Stroke	1 px borders. Prefer border over shadow.
Shadow	Almost none. If required: tight, dark, 0-blur.
Density	High on Flight Recorder. Marketing may breathe. Still severe.
Motion	120–200 ms ease. Instant preferred. No bounce, no confetti.
Veto flash	The only drama — short opacity pulse on danger-deep.
Icons	Stroke 1.5–2 px. Shield-off, gauge, ledger, waveform, lock.
Imagery	Product screens and line diagrams. No stock traders.

COMPONENT INVENTORY

Status pill	RUNNING · HALTED · KILL SWITCH — stamped, not cute.
Metric tile	Label + tabular value + delta.
Veto banner	Danger-deep field, mono numbers, reason + intermediates.
Dual gauge	Stated confidence vs calibrated p.
Equity chart	Guarded green / Shadow red, same axis, annotated DD.
Reliability	Stated x vs hit-rate y. Diagonal perfect-calibration line.
Decision table	Dense, expandable anomaly checklist.
Gate columns	G0 G1 G2 always side by side.
Check chips	S1 C1 M1 — small stamped IDs.
Sealed hash	limits_hash: a1b2... read-only treatment.

Flight Recorder first.

RUNNING EQUITY 10,482.40 PNL +1.4R OPEN 2 limits a1b2c9...

VETO — no_calibrated_edge p_adj 0.29 < breakeven 0.33 anomaly 0.12

REQUIRED PANELS

Overconfidence dial	Stated vs calibrated. Gap bars by confidence bucket.
Veto indicator	Reason + intermediates. Alarm that still reads as professional.
Equity curves	Guarded green vs Shadow red. Max drawdown annotated. Protect this.
Reliability diagram	Diagonal reference. Point size = n.
Decision log	Time, symbol, side, action, confidence, decision, p_adj, size.
Attack Lab	Scenario + attack picker. Three columns G0 G1 G2.
Results	HAR / ASR table. Reliability clean vs attacked.

LANDING PAGE ORDER (JUDGES)

Hook with numbers → Shared blind spots → How it works → Proof screens → Attack Lab A4 → Results → Repo / live / video

Do not default to Hero → six feature cards → pricing → FAQ. There is no pricing. Features are evidence panels.

STATES TO DESIGN

Empty · Running healthy · Veto just fired · Kill-switch / halted · Attack Lab replay

Empty copy: No proposals yet. Waiting for the next 15m close.

Numbers beat metaphors.

- 01 Every veto shows reason + intermediates (p_cal, p_adj, p_be, anomaly).
- 02 One definition of confidence everywhere: probability TP hits before SL.
- 03 Green / red semantic lock: guarded vs unguarded / harm. Never swap.
- 04 G0 / G1 / G2 always comparable side-by-side in Attack Lab.
- 05 If a metric looks bad, the UI still shows it.
- 06 The LLM is not the visual hero. The Inspector / brakes / evidence are.
- 07 Limits are sacred. Risk config is read-only. Show the hash.

GATES

G0 Unguarded

Fixed 1% risk on takes. The shadow book.

G1 Calibration

Limits + calibration + sizing. Catches crude attacks.

G2 Full gate

G1 + anomaly layer. Closes the A4 hole.

ATTACKS · DEMO NARRATIVE

- | | | |
|-----------|----------------------------|---|
| A1 | Sentiment injection | Instruction-like headlines → max confidence |
| A2 | Fake consensus | Duplicate bullish flood from fresh sources |
| A3 | Candle forgery | Rewrite trader candles into a textbook setup |
| A4 | Confidence steering | Land in a well-calibrated bucket (~74) to beat G1 |
| A5 | Adaptive | Iterates variants — stretch goal |

A4 is built to beat calibration-only. G1 lets it through. G2 should catch it and show which checks fired. Attacks change Trader inputs only. Inspector feed stays clean.

Do. Don't.

DO

- Keep monochrome discipline.
- Spend color only on meaning.
- Make Attack Lab + green/red equity the memorable move.
- Use real product terms from this kit.
- Design dark-first. Light is secondary.
- Optimize for a 3-minute judge glance.
- Show the shadow book. It is the moral.
- Keep Streamlit green/red roles on restyle.

DON'T

- Rebrand as generic AI-crypto SaaS.
- Add blue / purple / gold “for energy.”
- Soften vetoes into friendly toasts.
- Hide bad calibration.
- Invent social trading, NFTs, mobile.
- Use scales-of-justice as the identity.
- Neon grids, glitch, glassmorphism.
- Add a ® or write TareAI.

IMPLEMENTATION

- Live dashboard today: Streamlit + Plotly at dashboard/app.py.
- This kit is framework-agnostic. Apply tokens via CSS or Streamlit theme.
- Data: SQLite flight recorder. Demo seed via scripts/seed_demo_db.py.
- Design against seeded demo data + Attack Lab JSON — not live exchange DNS.
- Assets: logos/ for mark & lockups, tokens/tokens.css and tokens.json.

ABOUT · 50 WORDS

Tare is an agentic paper trader with a non-LLM Inspector. Setups come from market structure; an LLM proposes take/skip and a confidence. The Inspector sizes or vetoes using calibration, independent data, and hard limits. A shadow book and attack lab prove what the brakes catch — and what they miss.

Ship the brakes.

If a stranger can answer yes to all five, the kit is doing its job.

01 In 5 seconds: this is about distrusting AI confidence, not celebrating it.

02 In 15 seconds: they can point to stated vs calibrated, and green vs red equity.

03 In 60 seconds: they understand why a second LLM is the wrong fix.

04 In the Attack Lab: they see A4 pass G1 and fail G2 without a paper.

05 Nothing looks like default Streamlit, Binance promo, or Midjourney “AI trading.”



tare

BRAND KIT 1.1 · FORENSIC TRADING TERMINAL